

Experiments 1A–1C: Detailed Technical Report

Mathematical Gate, Discovery Atlas, and Adaptive Refinement
Query–Key Initialization Geometry Research Program

Automated report from frozen artifacts in `code/exp1a/data/full`, `code/exp1b/data/full`, and `code/exp1c/data/full`.

September 6, 2026

Contents

1	Executive Summary	3
2	Scientific Context and Motivation	4
2.1	What problem is being studied?	4
2.2	Why staged experiments before training a large LM?	4
3	Parameterization and Terminology Glossary	5
3.1	Coefficient tuple $(s, \alpha, \beta, \gamma)$	5
3.2	Token geometry notation	5
3.3	Attention-distribution metrics (1B)	5
3.4	Preregistered descriptive labels	6
4	Experiment 1A: Mathematical and Implementation Gate	7
4.1	Purpose	7
4.2	Experimental design	7
4.2.1	Configuration grid	7
4.2.2	Monte Carlo protocol	7
4.2.3	Hypothesis tests	8
4.2.4	Auxiliary checks (<code>checks.json</code>)	8
4.3	Results	8
4.3.1	Gate decision	8
4.3.2	Moment accuracy	8
4.3.3	Auxiliary check highlights	9
4.3.4	Interpretation of 1A	9
5	Experiment 1B: Discovery Atlas	10
5.1	Purpose	10
5.2	Reference geometry and grid	10
5.3	Measurement strata	10
5.4	Analysis pipeline	10
5.5	Pilot thresholds	11
6	Experiment 1B Results	12
6.1	Global label histogram	12
6.2	Synthetic vs. real activations	12
6.3	Reference stratum deep dive	12
6.4	Logit moments inside the 1B forward pass	13
6.5	Phase-map visualization	15
6.6	Label stability under threshold perturbation	18

7	Experiment 1C: Adaptive Refinement and Transfer	19
7.1	Purpose	19
7.2	Compute that was merged	19
7.3	Surrogate held-out metrics	19
7.4	Adaptive-cell labels	20
7.5	Why the portfolio did not freeze	20
7.6	Label sensitivity	20
8	Synthesis and Scientific Reading	21
8.1	What 1A established	21
8.2	What 1B established	21
8.3	What 1C established	21
8.4	What is <i>not</i> yet established	21
9	Limitations and Provenance	23
10	Recommended Next Steps	24
A	Holm–Bonferroni reminder	25
B	Bootstrap procedure in 1B	26

Chapter 1

Executive Summary

This document records the confirmatory analysis for the first three stages of the initialization-atlas program:

- **Experiment 1A** (mathematical and implementation gate): **PASSED**. All 864 Monte Carlo moment configurations and all auxiliary integration checks succeeded under preregistered tolerances.
- **Experiment 1B** (discovery atlas): **COMPLETE**. The full 1,717-cell coefficient grid was measured at reference geometry $(d, m, n) = (64, 16, 64)$ across 879,104 raw forward/backward records; bootstrap analysis produced 13,736 stratified summaries.
- **Experiment 1C** (adaptive refinement + transfer): **MEASURED, PORTFOLIO NOT FROZEN**. 120 surrogate-chosen cells $\times$ 2 geometries = 240 tasks, 122,880 raw rows, 1,920 stratified summaries. Held-out surrogate accuracy is high, but no coefficient point is `screened_candidate` on every replicated stratum, so `candidate_portfolio.json` remains unfrozen.

Headline 1B findings (all strata pooled, $N = 13,736$ bootstrap-labeled summaries):

Label	Count	Fraction
Diffuse	5,362	39.0%
Concentrated	3,793	27.6%
Screened candidate	2,867	20.9%
Boundary / uncertain	1,241	9.0%
Self-locked	472	3.4%
Gradient-starved	1	< 0.01%

These labels are **descriptive** statements about initialization-time attention geometry. They are *not* claims about downstream language-model quality until validated in Experiment 1D.

Headline 1C findings ($N = 1,920$ adaptive/transfer summaries):

Label	Count	Fraction
Diffuse	793	41.3%
Boundary / uncertain	613	31.9%
Screened candidate	514	26.8%
Concentrated / self-locked / gradient-starved	0	0%

The adaptive sampler targeted uncertain and theory-disagreement cells, so concentrated/self-locked pathology is absent by design. Transfer majority-label agreement between $(64, 16, 64)$ and $(128, 16, 128)$ is only 30/120 coefficient points.

Chapter 2

Scientific Context and Motivation

2.1 What problem is being studied?

Transformer attention computes, for each query position i and key position j ,

$$\ell_{ij} = \frac{1}{\sqrt{m}} (W_q x_i)^\top (W_k x_j), \quad P_{ij} = \frac{\exp(\ell_{ij})}{\sum_{j'} \exp(\ell_{ij'})},$$

where $W_q, W_k \in \mathbb{R}^{m \times d}$ are query and key projections and $x_i, x_j \in \mathbb{R}^d$ are token representations.

Standard initializers (e.g. Xavier/Kaiming on independent Gaussian weights) fix marginal variances but do not explicitly control the **joint geometry** of (W_q, W_k) . The research program introduces a **correlated Gaussian–orthogonal** family parameterized by $(s, \alpha, \beta, \gamma)$ (and angle ϕ), derives **exact finite-width** expressions for logit moments, and then **maps** where different attention behaviors occur in coefficient space before testing trainability and real-model effects.

2.2 Why staged experiments before training a large LM?

1. **1A** answers: “Is the theory implemented correctly?” If not, every map is meaningless.
2. **1B** answers: “Where in (s, α, ϕ) do diffuse, concentrated, self-locked, and screened behaviors occur at initialization?” This is an empirical atlas, not a hand-drawn phase diagram.
3. **1C** answers: “Do uncertain/boundary cells remain screened under a fitted surrogate, extra samples, and a width/length transfer?” A frozen portfolio is required before 1D.

Failure of 1A would block all downstream confirmatory interpretation. **1A passed.** Therefore 1B–1C maps may be interpreted as proposal evidence for the discovery atlas (subject to the limitations in Chapter 9). The 1C freeze rule has **not** yet produced a training portfolio.

Chapter 3

Parameterization and Terminology Glossary

3.1 Coefficient tuple $(s, \alpha, \beta, \gamma)$

Each admissible initialization is encoded by:

- $s > 0$: **scale** controlling overall query/key energy. We report $\log_2 s$ on atlas plots.
- $\alpha \in [-1, 1]$: **entrywise correlation** between flattened query and key coordinates within a head (target marginal correlation of Q_{ab} and K_{ab}).
- β, γ : split between a **row-space orthogonal** component B and an **independent Gaussian** component C in the construction $W_q W_k^\top = s^2(AB^\top + C)$ with constraints ensuring valid covariance.
- $\phi \in [0, \pi/2]$: angle parameterizing (β, γ) on the admissible circle at fixed α .

At $\alpha = \pm 1$, ϕ is redundant (all angles collapse to the same point); the grid retains one endpoint representative per sign.

3.2 Token geometry notation

For tokens x_i, x_j with Gram entries $v_i = \|x_i\|_2^2/d$, $v_j = \|x_j\|_2^2/d$, $c_{ij} = x_i^\top x_j/d$:

- **Predicted logit mean/variance**: closed-form $\mathbb{E}[\ell_{ij}]$ and $\text{Var}(\ell_{ij})$ from Lemma 1 (verified in 1A).
- $\hat{\tau}$: row-centered logit scale (empirical std. dev. of active logits in a row).
- $\hat{\Delta}$: **diagonal gap**—mean diagonal logit minus mean off-diagonal logit on the active set.

3.3 Attention-distribution metrics (1B)

For each attention row P_i over active keys $k_i > 1$:

Normalized entropy H_{norm}

$H(P_i)/\log k_i$. Equals 1 for uniform attention; near 0 for nearly one-hot rows.

Distance from uniform d_{unif}

RMS $\|P_i - u_i\|_2$ where u_i is uniform on the active support.

Maximum mass $\max_j P_{ij}$

Largest probability in the row (peakedness).

Self-mass P_{ii}

Diagonal attention probability (self-locking diagnostic).

Attention effective rank

$\text{erank}(P)/n$ where $\text{erank}(P) = \exp(-\sum_k \lambda_k \log \lambda_k)$ on singular values of the probability matrix.

Jacobian operator norm $\|J\|_2$

Spectral norm of softmax Jacobian $J(P_{i\cdot}) = \text{diag}(P_{i\cdot}) - P_{i\cdot}P_{i\cdot}^\top$ (local sensitivity / gradient health).

Update-to-weight ratio

$\|\nabla W\| \cdot \eta / \|W\|$ from a single backward step on a synthetic one-batch MSE probe (not full LM training).

Kernel metrics

Singular values of $M = W_q^\top W_k$ (operator norm, effective rank, spectral concentration).

3.4 Preregistered descriptive labels

Labels are assigned from **95% bootstrap confidence intervals** (2,000 seed-resampling draws per stratum in this run) with precedence:

self-locked $\succ$ concentrated $\succ$ diffuse $\succ$ gradient-starved $\succ$ screened candidate $\succ$ boundary/uncertain.

Diffuse

Lower CI bound on median $H_{\text{norm}} \geq 0.95$ and upper CI bound on $d_{\text{unif}} \leq 0.10$.

Concentrated

Upper CI bound on median $H_{\text{norm}} \leq 0.50$, *or* lower CI bound on median $\max_j P_{ij} \geq 0.80$.

Self-locked

Lower CI bound on median self-mass ≥ 0.80 .

Gradient-starved

Upper CI bound on Jacobian or update-to-weight below pilot floors (see Section 5.5).

Screened candidate

No pathology above confirmed; H_{norm} between 0.60 and 0.95; attention effective rank above pilot floor.

Boundary / uncertain

Intervals cross thresholds (genuine measurement uncertainty).

Important: “Screened candidate” means “passes forward pathology screens,” not “will train well.” Trainability is Experiment 1D.

Chapter 4

Experiment 1A: Mathematical and Implementation Gate

4.1 Purpose

1A is a **blocking quality-control stage**. It does not discover scientific novelty; it verifies that the codebase correctly implements:

- exact logit mean and variance formulas (Lemma 1);
- energy identities and entrywise variance targets;
- gradient/backprop identities (Proposition 4);
- common-bias cross-head statistic (Proposition 5);
- model integration (custom initializer, checkpoint survival, forward finiteness).

4.2 Experimental design

4.2.1 Configuration grid

- Widths $d \in \{32, 64, 128\}$.
- Head ratios $m/d \in \{1/8, 1/4, 1/2\}$.
- Token-pair correlations $c_{ij} \in \{-0.5, 0, 0.5, 1.0\}$ (fixed synthetic pairs).
- 24 coefficient points per geometry: four canonical corners plus 20 interior (α, ϕ) points.

Total moment configurations: $3 \times 3 \times 4 \times 24 = 864$.

4.2.2 Monte Carlo protocol

For each configuration:

1. Draw W_q, W_k from the admissible construction at the specified $(s, \alpha, \beta, \gamma)$.
2. Compute empirical mean and variance of $\ell = (W_q x)^\top (W_k y) / \sqrt{m}$ over increasing draws (minimum 20,000, up to 200,000).
3. Stop when $\text{SE}(\hat{\mu}) \leq 0.01 \sigma_{\text{pred}}$ and $\text{SE}(\widehat{\text{Var}}) \leq 0.01 \text{Var}_{\text{pred}}$.
4. Record energies, orthogonality $\|AB^\top\|_F / (\|A\|_F \|B\|_F)$, entry variances, empirical entry correlation vs. α .

4.2.3 Hypothesis tests

For each configuration we form two-sided normal approximate p -values for mean and variance. Across all 864 configurations (1,728 tests), **Holm–Bonferroni** correction controls family-wise error at $\alpha = 0.01$. A configuration **passes** only if:

- neither corrected mean nor variance test is rejected;
- precision targets met;
- relative orthogonality $< 10^{-6}$;
- component energies within 1% of s^2m (and B, C within 1% of m);
- entry variances within 2% of s^2/d ;
- $|\hat{\rho}_{\text{entry}} - \alpha| \leq 0.02$.

4.2.4 Auxiliary checks (checks.json)

- **Gradient identity**: 20 random small configurations; analytic vs. finite-difference relative error $< 10^{-5}$.
- **Common-bias ratio**: empirical vs. predicted cross-kernel correlation ratio; error < 0.02 .
- **Baseline initializer match**: KS test comparing custom independent point to i.i.d. Gaussian baseline ($p > 0.01$, variance ratio in $[0.9, 1.1]$).
- **Integration**: initializer changes W_q, W_k ; checkpoint reload is bitwise-equal; first forward pass finite.

4.3 Results

4.3.1 Gate decision

gate_summary.json: passed = true

Moment configurations tested: 864 Moment failures: 0 Gradient failures: 0

4.3.2 Moment accuracy

- All 864 configurations achieved **precision_converged = true**.
- Holm-rejected mean tests: **0 / 864**. Holm-rejected variance tests: **0 / 864**.
- Maximum relative orthogonality across configurations: $\lesssim 10^{-16}$ (numerical zero).
- Uncorrected p -values span $(3 \times 10^{-4}, 1)$, consistent with valid null sampling rather than systematic bias.

Figure 4.1 is the actual moment-validation plot: empirical vs. predicted with a $y = x$ line, plus residuals. Calibration regression on the 864 configs is

$$\hat{\mu} = -3.4 \times 10^{-5} + 0.9999 \mu_{\text{pred}}, \quad R^2 = 0.99999, \quad \text{RMSE} = 0.0058,$$

$$\widehat{\text{Var}} = 0.0024 + 0.9976 \text{Var}_{\text{pred}}, \quad R^2 = 0.9975, \quad \text{RMSE} = 0.0106.$$

Median relative variance error is 0.62% (95th percentile 1.8%). This—not the 1B sequence scatter—is the evidence that Lemma 1 is implemented correctly.




Figure 4.1: Experiment 1A: isolated-pair Monte Carlo moments versus exact theory. Residuals are centered at zero with magnitude $\sim 10^{-2}$, matching the preregistered precision gate.

4.3.3 Auxiliary check highlights

- Common-bias absolute error: 1.76×10^{-4} (threshold 0.02).
- Baseline KS $p = 0.574$, variance ratio 1.009.
- All 20 gradient configurations: identity error 0; finite-difference relative errors $\sim 10^{-10}$.
- Integration check: initializer changed tensors, checkpoint bitwise equal, forward finite.

4.3.4 Interpretation of 1A

The implementation faithfully realizes the finite-width theory within Monte Carlo tolerance. **Downstream atlas interpretation is unblocked.** Any future failure in 1B-Exp2 cannot be attributed to a Lemma 1 coding error without revisiting this gate.

Chapter 5

Experiment 1B: Discovery Atlas

5.1 Purpose

1B empirically maps **initialization-time attention behavior** across the admissible coefficient domain. Unlike 1A’s isolated token pairs, 1B runs a full forward pass through a one-block attention head on **sequences** of length $n = 64$ with value projection and (for the backward probe) a single optimization step.

5.2 Reference geometry and grid

Parameter	Value
Token dimension d	64
Head dimension m	16
Sequence length n	64
$\log_2 s$ grid	17 values in $[-2, 2]$
α grid	13 values in $[-1, 1]$
ϕ grid	9 values in $[0, \pi/2]$
Unique cells	1,717
Seeds per cell	64 paired seeds

5.3 Measurement strata

Each cell is evaluated under **eight** input/mask strata (per ϕ slice, aggregated in analysis):

1. Synthetic inputs with token correlation $r = 0, 0.25, 0.5$ (LayerNorm applied).
2. **Real activations**: post-LayerNorm tensors from an independently initialized baseline decoder (distribution shift test).
3. **Causal** (lower-triangular) vs. **unmasked** (full) attention.

Raw records per full run: $1,717 \times 64 \times 8 = 879,104$ rows (matches `analysis_summary.json`).

5.4 Analysis pipeline

1. Load all `atlas_shard*.jsonl` shards (~ 1.3 GB JSONL).
2. Compute pilot floors from independent baseline $(s, \alpha, \gamma) = (1, 0, 1)$ (Section 5.5).
3. Group rows by $(\text{cell_id}, \text{input_kind}, r, \text{mask}, d, m, n)$.

4. For each group and each of seven metrics, compute median and **cluster bootstrap 95% CI** with 2,000 draws resampling the 64 seeds.
5. Assign label + ± 0.05 sensitivity labels.
6. Export `atlas_summary.csv` (13,736 rows) and 72 stratum phase maps. The analyze-script residual scatter is *not* a 1A replication (Section 6.4).

5.5 Pilot thresholds

Derived from the independent baseline at $s = 1$ (5th percentile across baseline records):

Floor	Value
Jacobian operator norm	0.0556
Update-to-weight (min of W_q, W_k)	9.75×10^{-6}
Attention effective rank	0.275

Gradient-starved labels are rare when floors are low; only **one** stratum in the entire atlas was labeled gradient-starved.

Chapter 6

Experiment 1B Results

6.1 Global label histogram

Pooling all 13,736 stratified summaries:

Label	Count	Percent	Qualitative reading
Diffuse	5,362	39.0%	Near-uniform attention rows
Concentrated	3,793	27.6%	Low entropy or high peak mass
Screened candidate	2,867	20.9%	Mid-entropy, passes pathology screens
Boundary / uncertain	1,241	9.0%	Thresholds not resolved at 95%
Self-locked	472	3.4%	High diagonal self-attention
Gradient-starved	1	< 0.01%	Jacobian/update below pilot floor

6.2 Synthetic vs. real activations

Input kind	Diffuse	Concentrated	Screened	Boundary	Self-locked	Grad-starved
Synthetic	4,115	2,811	2,115	960	300	1
Real activation	1,247	982	752	281	172	0

Screened fraction: **24.3%** synthetic vs. **21.5%** real (similar order of magnitude; real activations slightly shift mass toward concentrated/self-locked regions).

6.3 Reference stratum deep dive

We highlight **synthetic, unmasked**, $r = 0.25$, $\phi = 0$ —the canonical “paper slice” with 221 coefficient points (one per admissible (s, α) at this ϕ).

Label	n	Mean H_{norm}	Mean max mass	Mean self-mass	Role
Diffuse	99	0.990	0.028	0.017	High-entropy plateau
Screened candidate	40	0.829	0.141	0.038	Candidate band
Concentrated	60	0.186	0.761	0.076	Peaked attention
Self-locked	8	0.022	0.980	0.971	Diagonal collapse
Boundary	14	0.613	0.301	0.044	Uncertain

At $\phi = \pi/4$ on the same stratum ($n = 187$ cells), counts are similar (85 diffuse, 33 screened, 53 concentrated, 3 self-locked), indicating ϕ rotates mass between orthogonal/independent components without collapsing the map.

6.4 Logit moments inside the 1B forward pass

1A tests pairwise moments with many i.i.d. draws of (W_q, W_k) . 1B stores a *different* diagnostic: for each seed it compares one realized $n \times n$ logit matrix to the closed-form gram formulae. These quantities are not interchangeable.

What is stored. Empirical `logit_mean` / `logit_variance` are the sample mean and sample variance of the realized masked logits $\{\ell_{ij}\}$. Predicted `predicted_logit_mean` is the average of the pairwise conditional means $\mathbb{E}[\ell_{ij} | x]$. Predicted `predicted_logit_variance` is the average of the pairwise *conditional* variances $\text{Var}(\ell_{ij} | x)$, **not** the across-pair variance of the map. By the law of total variance,

$$\text{Var}_{ij}(\ell_{ij}) = \mathbb{E}_{ij}[\text{Var}(\ell_{ij} | x)] + \text{Var}_{ij}(\mathbb{E}[\ell_{ij} | x]).$$

The second term is large whenever αs^2 is large (token Gram heterogeneity). Plotting empirical across-pair variance minus mean pairwise conditional variance therefore *must* show structured bands at discrete s values. That is an estimand mismatch, not a 1A failure.

The retracted statistic. The previous draft quoted mean residuals -0.0048 ± 0.0484 and $+0.0003 \pm 0.0013$ as “excellent agreement.” Those numbers used only the first 20,000 raw rows (cell ids 0–39, the smallest-scale cells). The analyze-script figure itself was drawn on all 879,104 records. The two do not match, and the global mean residual can cancel opposite stratum biases (e.g. $r = 0$ vs. $r = 0.5$).

Full-data calibration (this report). On all 879,104 single-draw records, OLS of empirical vs. predicted is:

Estimand	a	b	R^2	RMSE	Median rel. err.
Mean (raw records)	0.004	1.001	0.788	1.77	4.7%
Mean (seed-averaged, 13,736 strata)	0.004	1.001	0.996	0.21	0.64%
Variance, pairwise formula (raw)	0.14	1.026	0.920	20.5	17%
Variance, pairwise (seed-averaged)	0.13	1.027	0.949	16.1	16%
Variance, total-variance correction (raw)	−0.06	0.975	0.936	18.3	14%
Variance, total-variance (seed-averaged)	−0.07	0.975	0.965	13.3	8.0%

Mean prediction is calibrated ($b \approx 1$). Single-draw residuals are wide (RMSE = 1.77); averaging 64 seeds recovers $R^2 = 0.996$. That supports the mean formula at sequence scale after Monte Carlo averaging, not per-record precision.

Variance is **not** “excellent agreement.” The pairwise residual mean on the full file is $+0.96 \pm 20.4$, not $+0.0003$. Discrete islands in the residual plot coincide with the 17 values of s . Stratified pairwise variance residuals (raw records):

Stratum	Mean residual	RMSE	OLS slope
Synthetic inputs	−0.55	20.4	0.98
Real activations	+5.46	20.7	1.18
Unmasked	+0.27	18.8	1.00
Causal	+1.64	22.0	1.05
$r = 0$	+5.37	20.5	1.18
$r = 0.25$	+0.96	15.3	1.03
$r = 0.5$	−7.97	24.3	0.75

Positive and negative stratum biases cancel in the pooled mean. Adding an off-diagonal Gram estimate of $\text{Var}_{ij}(\mathbb{E}[\ell_{ij} \mid x])$ improves R^2 and centers residuals by s , but heteroscedasticity remains (spread grows with s). The 1A formula is still the right pairwise theory; the 1B stored variance is a different functional of the sequence map.




Figure 6.1: 1B single-draw diagnostics. Left: mean is approximately unbiased with heteroscedastic scatter. Right: empirical across-pair variance vs. mean pairwise conditional variance—mismatched estimands, discrete s -bands, large residuals. Color is $\log_2 s$.




Figure 6.2: Same empirical variance after adding $\text{Var}_{ij}(\mathbb{E}[\ell_{ij} \mid x])$. Residuals are closer to centered by scale; remaining error is heteroscedastic single-draw noise, not 1A-style 10^{-2} Monte Carlo error.




Figure 6.3: Variance residuals boxed by s . Naive pairwise residuals fan out and shift with scale; the total-variance correction recenters medians. Whiskers at $s = 4$ still span $\mathcal{O}(10^2)$.

Conclusion for this section: 1A remains the confirmatory moment test. 1B means track the theory after seed averaging. 1B variances require matching the across-pair estimand and stratified reporting; they do not license a claim of excellent per-record sequence-scale variance agreement.

6.5 Phase-map visualization

Figure 6.4 follows the reference style used in information-theory phase diagrams: hatched region legend, dashed operational contours, and red-edged dots on the measured $(\log_2 s, \alpha)$ grid. Region C (screened candidate) is drawn as a **narrow** horizontal band to avoid overstating its spatial extent.




Figure 6.4: Region-style empirical atlas for the reference stratum. Dashed curves trace interpolated median H_{norm} and self-mass fields; colored dots are measured cells with preregistered labels.

Figure 6.5 shows the same data as a smooth surface (analogous to a 3D capacity/entropy landscape).

Normalized row-entropy surface over coefficient space




Figure 6.5: Interpolated normalized row-entropy surface over $(\log_2 s, \alpha)$.

Figure 6.6 gives the discrete nearest-cell label mosaic.




Figure 6.6: Nearest-measured-cell label assignment on the coefficient grid.

Legacy scatter plot (analysis script default):




Figure 6.7: Original turbo-colormap scatter from `exp1b/run.py` analyze.

Additional 71 figures (varying ϕ , mask, r , and real activations) are in `code/exp1b/data/full/figures/`.

6.6 Label stability under threshold perturbation

Comparing primary labels to ± 0.05 probability/entropy shifts:

- 1,294 strata (9.4%) change label under -0.05 shift.
- 6,721 strata (48.9%) change under $+0.05$ shift.

Nearly half the atlas lies near label **boundaries** under generous threshold perturbation. This is expected: softmax metrics transition continuously. Stage 1C adaptive refinement was run to probe these boundaries (next chapter).

Chapter 7

Experiment 1C: Adaptive Refinement and Transfer

7.1 Purpose

1C fits a probabilistic surrogate on 1B discovery labels, holds out 20% of coefficient cells, proposes 120 uncertain / theory-disagreement points, and remeasures them at

$$(d, m, n) \in \{(64, 16, 64), (128, 16, 128)\}$$

with 64 seeds and the same eight input/mask strata as 1B. The freeze rule then requires a coefficient $(s, \alpha, \beta, \gamma)$ to be labeled `screened_candidate` on *every* replicated stratum before it can enter `candidate_portfolio.json`.

7.2 Compute that was merged

Kaggle produced two archives (`exp1c.tar`, `exp1cccc.tar`). Incomplete shards 0002 and 0003 from the first archive (5/30 tasks) were replaced by the completed copies. Effective sharding was 8 GPU streams covering all 240 tasks.

Item	Value
Adaptive cells	120
Geometries	2
Tasks	240 (all present)
Raw rows	122,880 (512 per task)
Stratified summaries	1,920
Bootstrap draws	2,000
Pilot floors	Jacobian 0.0241, update 5.94×10^{-6} , rank 0.175

7.3 Surrogate held-out metrics

The classifier never trains on the held-out 20% of 1B cells. Reported metrics:

Metric	Value
Train / test rows	10,984 / 2,752
Balanced accuracy	0.987
Multiclass log-loss	0.0355
Expected calibration error	0.0153
Boundary error (top-quartile uncertainty)	0.0349

The interpolator is well calibrated on 1B labels. It is **not** itself evidence that those labels are trainable.

7.4 Adaptive-cell labels

Label	All 1C	(64, 16, 64)	(128, 16, 128)
Diffuse	793	266	527
Boundary / uncertain	613	314	299
Screened candidate	514	380	134
Other pathologies	0	0	0

The transfer geometry is substantially more diffuse and less often screened. Majority-label agreement across the two geometries is 30/120 coefficient points.

On the canonical 1B slice (synthetic, unmasked, $r = 0.25$, reference geometry) the 120 adaptive points are 94 diffuse, 25 boundary, and 1 screened: the sampler did what it was asked to do—it spent budget on uncertain cells, not on the high-entropy plateau.

7.5 Why the portfolio did not freeze

`freeze_portfolio` groups rows by $(s, \alpha, \beta, \gamma)$ and keeps a point only if **all** 16 copies (2 geometries $\times$ 2 masks $\times$ 4 input correlations/kinds) are `screened_candidate`.

Criterion (120 coefficient points)	Count
Any stratum screened	117
Majority of strata screened	5
All 8 reference-geometry strata screened	0
All 16 replicated strata screened (freeze rule)	0

`candidate_portfolio.json` therefore records `frozen: false` (“need 8 replicated screened candidates; found 0”). This is a **replication-strictness** outcome, not a failed measurement. 1D cannot start on a frozen portfolio until the rule is relaxed (e.g. reference geometry only, or majority vote) or more cells are measured.

7.6 Label sensitivity

Of 1,920 1C strata, 534 (27.8%) change label under a -0.05 probability/entropy shift and 1,404 (73.1%) change under $+0.05$. Adaptive cells sit even closer to decision boundaries than the 1B atlas, which is expected.

Chapter 8

Synthesis and Scientific Reading

8.1 What 1A established

The finite-width algebraic theory is not merely analytic—it is implemented correctly in floating point with rigorous Monte Carlo and gradient checks. This is a prerequisite most initialization papers skip.

8.2 What 1B established

1. The admissible (s, α, ϕ) domain admits **large, structured regions** of distinct initialization phenotypes (diffuse vs. concentrated vs. screened).
2. **Self-locking** is measurable but occupies a minority ($\sim 3.4\%$ of strata).
3. **Gradient starvation** at initialization is extremely rare under the pilot floors (1 stratum).
4. Pairwise mean identities carry over to 1B after seed averaging ($R^2 = 0.996$). Sequence-scale *variance* as stored in 1B is a different estimand and is only approximately calibrated after a total-variance correction (Section 6.4); it does not re-prove 1A.
5. Real activations preserve qualitative label proportions but redistribute mass toward concentration relative to synthetic Gaussians.

8.3 What 1C established

1. A discovery surrogate can be fit and held out with balanced accuracy 0.987 and ECE 0.015.
2. The 120 adaptive cells were fully remeasured at two geometries (122,880 rows).
3. Screened labels are **not** stable under mask, correlation, real activations, and width/length transfer jointly: the preregistered freeze intersection is empty.
4. Transfer to (128, 16, 128) shifts mass toward diffuse (527 vs. 266 at reference).

8.4 What is *not* yet established

- A frozen 8-head `candidate_portfolio.json` under the current all-stratum rule.
- Whether screened candidates **train better** (Experiment 1D).
- Whether a diverse portfolio improves a 32M LM (Experiment 2).

- Sign-symmetry unit test on 50 cells (`sign_symmetry_check.json` not present in the full 1B analyze artifact).

Chapter 9

Limitations and Provenance

- **Single-block probe:** no depth, no residual stream, no full LM in 1B.
- **One backward step:** update ratios are diagnostic, not training trajectories.
- **Bootstrap draws:** 2,000 (manifest `exp1b_1788602938369_18144.json`); coarser than pre-registered 10,000 but sufficient for exploration.
- **CPU/RAM analyze:** ~ 1.3 GB JSONL loaded fully into memory.
- **Labels are descriptive**, not causal claims about generalization.
- **1B logit variance is not the 1A estimand:** `predicted_logit_variance` averages pairwise conditional variances; `logit_variance` is the across-pair sample variance of one realized map. Pooled mean residuals cancel opposing r and input-kind biases.
- **1C freeze is strict:** a point must be screened on all 16 replicated strata. Zero of 120 points pass; 1D is blocked until the rule is changed or more cells are collected.
- **Git LFS:** raw shards require `git lfs pull` or manual upload on cloud machines.

Artifact paths:

- 1A: `code/exp1a/data/full/gate_summary.json`, `moment_results.json`, `checks.json`
- 1B: `code/exp1b/data/full/analysis_summary.json`, `atlas_summary.csv`, `figures/`
- 1C: `code/exp1c/data/full/analysis_summary.json`, `atlas_summary.csv`, `candidate_portfolio.json`, `held_out_metrics.json`

Chapter 10

Recommended Next Steps

1. **Unblock 1C freeze:** either require screening only on the reference geometry / majority vote, or propose additional cells around the 5 majority-screened coefficients.
2. **1D:** once a portfolio is frozen, train one-block decoder on associative recall + TinyStories; test screened vs. pathology early-loss gate.
3. **Exp2:** 32M LM confirmatory study with frozen treatments.

Appendix A

Holm–Bonferroni reminder

Given ordered p -values $p_{(1)} \leq \dots \leq p_{(m)}$, reject hypothesis i if $p_{(i)} \leq \alpha/(m - i + 1)$. We applied $m = 1728$, $\alpha = 0.01$ in 1A.

Appendix B

Bootstrap procedure in 1B

For each stratum with seed values $\{x_s\}_{s=1}^{64}$:

1. Resample 64 seeds with replacement $B = 2000$ times.
2. Compute median metric per resample.
3. Report 2.5% and 97.5% quantiles as CI endpoints.

Seeds are paired across cells; heads/tokens nested within seed are not separate bootstrap units (cluster bootstrap at seed level).